

CSC311 Notes

When to use machine learning vs. other computer tools?

- It's hard to manually code a solution to an issue (eg. recognizing vision, speech)
- Program has to adapt to changing environments (eg. detect spam)
- Program should perform better than humans
- Privacy/fairness (eg. ranking search results)

Machine learning vs statistics

- Both find patterns in data, draw on similar math bases, same core algorithms
- Stats about helping scientists/policymakers draw conclusions; ML about algorithms, coding the theoretical
- Stats about interpreting results, math & rigor; ML about predictive performance scalability, autonomous agents

Programs learn from experience E with respect to tasks T , performance measures P , when E grows $\Rightarrow P$ improves.

Supervised Learning: Teaching by giving labelled examples of correct behavior.

Reinforcement Learning: Teaching by letting program interact with simulation, maximize some value

Unsupervised Learning: No labelled examples, finds “interesting” data patterns.

Supervised Learning

Parameter: A variable that affects the learning model's output. “Learning” refers to optimizing the parameters.

Parametric Models: Models that “learn” parameters from data and at test-time, refer back to the “learned” data

- eg. linear regression, which learns optimal weights for a linear equation like $y(x) = \alpha x + \beta$
- eg. decision trees, which build a tree based on parameter values for predicting stuff

Non-Parametric Models: Models without parameters. There's no “learning phase”; all work is done at run-time.

- eg. nearest neighbours

Training Set (\mathcal{D}): A set of data fed into the program for learning.

- Data consists of examples of correct behavior to emulate: **inputs** and **labels** (i.e. correct outputs)
- Labels can also be called targets/responses/outcomes/outputs/classes
- Usually, we collapse any input into a **vector** of form $x = \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}$ because it makes computation fast.
 - The vector components are also called **features/covariates**
 - e.g. Images can be thought of huge matrices of RGB values. We concatenate each row into a super-long vector because we don't need the matrix's extra “spatial” information.
- **Regression Algorithm:** An algorithm $\mathbb{R}^d \rightarrow \mathbb{R}$ that predicts continuous values
- **Classification Algorithm:** An algorithm $\mathbb{R}^d \rightarrow \{\text{discrete outputs}\}$ that predicts a class
- Formally, $\mathcal{D} = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_n, t_n)\}$, where $\mathbf{x}_i \in \mathbb{R}^d$ is an input and $t_i \in \mathbb{R}$ is the corresponding label.

Validation Set: A set of data where for any (\mathbf{x}_i, t_i) pair, each t_i is compared to $y(\mathbf{x}_i)$ using the same model on different hyperparameter values. Used to find the optimal hyperparameter values.

- **Hyperparameter:** A parameter we set that affects learning, in contrast to parameters an algorithm derives.
- **Grid Search:** Exhaustively searching for optimal value through a manually-chosen subset of hyperparameters
- **Random Search:** Testing random hyperparameter configurations using a probability distribution

Test Set: A set of data where for any (\mathbf{x}_i, t_i) pair, t_i is compared to $y(\mathbf{x}_i)$, the model's prediction

- **Generalization Error:** How often the program correctly predicts data it hasn't seen before (ie. **generalizes**).
- Test sets are used on the “final” model, unlike the validation set, which tests variations of the same model

Nearest Neighbours

Argmin (argmin $f(x)$): The argument of $\min_{x \in \mathcal{D}} f(x)$; $f\left(\argmin_{x \in \mathcal{D}} f(x)\right) = \min_{x \in \mathcal{D}} f(x)$. Same logic with argmax.

Indicator/Identity Function: The function $\mathbb{I}(P) = \begin{cases} 1 & P \text{ is true} \\ 0 & P \text{ is false} \end{cases}$.


Metric Space: A set of points with some notion of distance between them defined by a distance function. A generalization of the idea of distance (MAT327).

L^p Space: In functional analysis, a metric space defined by distance function

$$\|\mathbf{x}\|_p = \sqrt[p]{\sum_{i=1}^d |x_i|^p} = \sqrt[p]{|x_1|^p + |x_2|^p + \dots + |x_n|^p}$$

where $\mathbf{x} \in \mathbb{R}^d$. This is also called the **p -norm**.

- On the right is the unit circle in \mathbb{R}^2 under different p -values.
- The **2-norm** is the Euclidean norm, how we've understood distance so far: $\|\mathbf{x}\|_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}$



Nearest Neighbours (NN): A non-parametric algorithm $y(\mathbf{x})$ that, given a query point \mathbf{x} , returns the label of the training set input “nearest” to \mathbf{x} .

- Since we assume inputs are vectors, use Euclidean distance: $\|\mathbf{u} - \mathbf{v}\|_2 = \sqrt{\sum_{i=1}^d (u_i - v_i)^2}$

Not always the *best* distance measure. Higher/lower p -values make larger/smaller errors more significant.

Algorithm

- Recall the training set $\mathcal{D} = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_n, t_n)\}$. Find the nearest (\mathbf{x}^*, t^*) to \mathbf{x} with the formula

$$\mathbf{x}^* = \argmin_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} (\text{distance}(\mathbf{x}_i, \mathbf{x})) = \argmin_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} \|\mathbf{x}_i - \mathbf{x}\|$$
 In other words, $\|\mathbf{x}^* - \mathbf{x}\| = \min_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} \|\mathbf{x}_i - \mathbf{x}\|$. If there are multiple minimums, pick any one of them.
- Output $y(\mathbf{x}^*) = t^*$

Practically, compute $\|\mathbf{x}_i - \mathbf{x}\|^2 = \sum_{i=1}^d (u_i - v_i)^2$ since it saves computing $\sqrt{\dots}$ and squaring preserves argmins.


- Algorithm is the same between regression/classification; only difference is t^* is a scalar/discrete value
- Despite its simplicity, NN is a competitive learning algorithm
- NN is very effective if we know *what* we should measure distances of, and *how* we measure distances

Input Space: Set of all possible inputs (vectors) that the learning model accepts; the domain of the learning model's algorithm, in \mathbb{R}^d

Decision Boundaries: The dividing lines in input space between inputs that output different categories.

Voronoi Diagram: Given points $\{p_1, \dots, p_n\}$ in \mathbb{R}^d , a partition of input space into n tiles where each tile contains the closest points to some p_i .

- Helps visualize how NN assigns parts of input space to categories
- Plot \mathbf{x} , see the tile \mathbf{x} falls in, find the other point \mathbf{x}^* in the tile, return t^*



K-Nearest Neighbours (KNN): An extension of NN using odd hyperparameter $k \in \mathbb{N}$. Finds the k nearest neighbours to X and returns a normal/weighted average (regression) or majority classification (classification).

Algorithm

- 1) Find k examples of (\mathbf{x}^*, t^*) with the closest \mathbf{x}^* values to \mathbf{x} using the NN algorithm.
- 2) Output the mean/mode of $\{t_1^*, \dots, t_k^*\}$ with the formulae

Regression (Mean)

$$y(\mathbf{x}) = \frac{1}{k} \sum_{i=1}^k t_i^*$$

Classification (Mode)



$$y(\bar{\mathbf{x}}) = \underset{t_i^* \in \{t_1^*, \dots, t_k^*\}}{\operatorname{argmax}} \sum_{i'=1}^k \mathbb{I}(t_i^* = t_{i'}^*)$$



Small k	Balanced k	Large k
<ul style="list-style-type: none"> • <u>Overfitting</u>: Too sensitive to noise or mislabelled data, complex ("class noise") • Captures fine-grained patterns • If $k = 1$, training set error is 0 	<ul style="list-style-type: none"> • Nice theoretical properties if $\lim_{n \rightarrow \infty} k = \infty$ and $\lim_{n \rightarrow \infty} \frac{k}{n} = 0$ with $k = f(n)$ for some function f • Generally, choose $f(n) < \sqrt{n}$ • Can be chosen by validation sets 	<ul style="list-style-type: none"> • <u>Underfitting</u>: Simple, can't show important regularities • Stable predictions by averaging examples • If $k = n$, KNN gives the same output for any input

Validation, Testing, and KNN

- Use a validation set to tune for k , then a testing set at the end to test everything
- As k grows, training error increases, but this is fine (don't expect 100% accuracy)
- What matters is performing just as well on training and testing sets; generalizing



Pitfalls of Nearest Neighbours

Curse of Dimensionality: Various phenomena with data in higher dimensions, occurring usually because *volume of space increases very fast* with increasing dimensions

- e.g. Sample points from n -dimensional sphere. As n increases, distance between points get **farther** and **more uniform** (*provable using expected value and covariance*)
- Thus, common distances become less meaningful; using Euclidean norm may become problematic.
- Luckily, some data has low **intrinsic dimension**: it lies on or near a low-dimensional manifold, so the volume issues are less pronounced
 - **Manifold:** Informally, a lower-dimensional surface/object in a higher dimension

Normalization: Issues regarding units – higher ranges of absolute numbers in a vector contributes more to distance than smaller ranges of numbers.

- e.g. In a dataset where $\mathbf{x} = \begin{bmatrix} \text{age} \\ \text{income} \end{bmatrix}$, age ranges from 0 to 100, income ranges from 0 to 500,000. Income will influence distance much more than age.
- Normalize each vector by setting the boundaries to 0 and 1 by applying this on each value:

$$x_i \leftarrow \frac{x_i - \mu}{\sigma}$$

where μ = mean of all x_i , σ = standard deviation of all x_i .

- However, there are also cases where we don't want to normalize!

Computational Cost: In machine learning, which requires large amounts of data, there are very high standards for efficient running-times of algorithms/memory storage

- Naïve implementations of NN have running-time $\mathcal{O}(nd)$, with n data points, d dimensions, which then must be sorted, which is usually $\mathcal{O}(n \log n)$. This algorithm is run for each query, too, which is expensive.
- A lot of work has been done to make NN more efficient.

Decision Trees

Entropy: In information theory, a number describing information/surprise/uncertainty in a random variable (RV).

- Symbol is H , capital letter of η
- Unit is a **bit**, because it's commonly used in computers storing/compressing/encrypting information
- **High Entropy:** Variable distribution is uniform; less predictable; more "info"
 - e.g. A fair coin flip has entropy of 1 bit
- **Low Entropy:** Variable distribution is concentrated, more predictable; less "info"; 0 means 100% predictability
 - e.g. A biased coin flip of 8/9 heads, 1/9 tails has an entropy of $\sim 1/2$ bits
- Most commonly applies to discrete RVs, but can be extended to continuous RVs in *differential entropy*

$$H(X) = - \sum_{x \in X} p_X(x) \log_2 p_X(x)$$

$$H[Y|X = x] = - \sum_{y \in Y} p_{Y|X}(y|x) \log_2 p_{Y|X}(y|x)$$

$$H(X, Y) = - \sum_{x \in X} \sum_{y \in Y} p_{X,Y}(x, y) \log_2 p_{X,Y}(x, y)$$


$$\mathbb{E}[H[Y|X]] = - \sum_{x \in X} \sum_{y \in Y} p_{X,Y}(x, y) \log_2 p_{Y|X}(y|x)$$

$$H(X) \geq 0 \quad \left| \quad H(X|X) = 0 \quad \left| \quad \begin{array}{l} H(X|Y) = H(X) \\ \text{for independent } X, Y \end{array} \right. \quad H(X, Y) = H(X|Y) + H(Y) \quad \left| \quad H(X|Y) \leq H(X) \right.$$

Information Gain ($IG(Y|X)$): How much more certain we are about Y 's result from knowing X .

$$IG(Y|X) = H(Y) - H(Y|X)$$

- AKA how informative X is to Y , or the loss of entropy/uncertainty in Y due to knowing X
- X is *completely informative*: $IG(Y|X) = H(Y)$
- X is *completely uninformative*: $IG(Y|X) = 0$



Greedy Heuristic: A type of approach/strategy that involves always choosing the most locally-optimal choice (ie. what looks best in the moment)

Loss (\mathcal{L}): Refers to a loss function, that returns a non-negative scalar quantifying "error levels". Should be minimized

- **Objective Function:** In optimization, a function that needs to be minimized.
- **Cost Function:** Usually averaged loss over \mathcal{D} plus the regularizer function (see later). No strict definition, often interchangeable with loss. (sometimes, omit $\frac{1}{n}$ as it's a scalar multiple)


$$\mathcal{J}(\mathcal{D}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(y(\mathbf{x}_i), t_i)$$

- **Mean Squared Error:** Average squared difference between a model output $y(\mathbf{x})$ and training target t .

$$\mathcal{J}(\mathcal{D}) = \frac{1}{n} \sum_{i=1}^n (y(\mathbf{x}_i) - t_i)^2, \quad \mathcal{L}(y(\mathbf{x}), t) = (y(\mathbf{x}) - t)^2$$

Decision Tree: A parametric model based on a tree structure where features are split based on their value.

- **Decision boundaries** are axis-aligned lines
- **Internal nodes** are booleans about feature values
- **Leaf nodes** are outputs (ie. predictions)
- **Hyperparameters** are tree height, branches per node, number of features, nodes, information gain threshold, etc.



Algorithm – Learning

- 1) Start with no nodes.
- 2) Use a **greedy heuristic** to pick features and a split that minimizes **loss**.
 - a. *Regression Tree*: Minimize **squared error** in the divided regions per split
 - b. *Classification Tree*: Maximize **information gain** per split

For splits with the same loss, the program picks splits that maximize the decision boundary's margins – i.e. distance to points on both sides

- 3) Stop subdividing if all points in regions have same classification, no points left, or hyperparameters say so

Algorithm – Predicting

- 4) Let R be a region of input space corresponding to an output. Find the R associated with new point \mathbf{x} .
Let $\{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_k, t_k)\}$ be the training points where $x_i \in R$
 - a. *Regression Tree*: Since the output is continuous, we output the mean of all t_i
 - b. *Classification Tree*: Since the output is discrete, we output the mode of all t_i

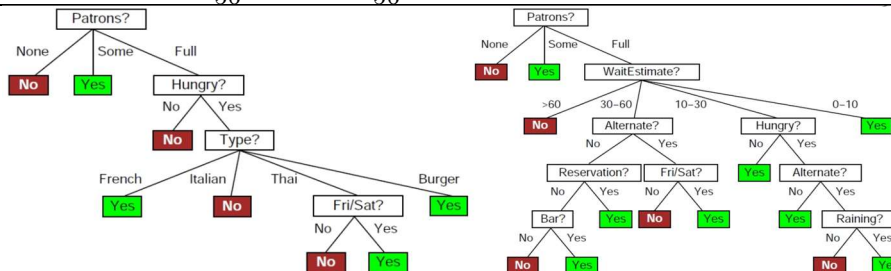
eg. A decision tree classifies * and *. What is the information gain of the split as shown on the right?

Original:	13 *, 56 *	Split:	Left: 3 *, 22 *
			Right: 10 *, 21 *

$$\begin{aligned}
 H(Y) &= - \sum_{y \in Y} p_Y(y) \log_2 p_Y(y) \\
 &= -p_Y(*) \log_2 p_Y(*) - p_Y(\star) \log_2 p_Y(\star) \\
 &= -\frac{13}{56} \log_2 \frac{13}{56} - \frac{43}{56} \log_2 \frac{43}{56} \\
 &\approx \boxed{0.7817}
 \end{aligned}$$

$$\begin{aligned}
 H(Y|X = \text{left}) &= - \sum_{y \in Y} p_{Y|X}(y|\text{left}) \log_2 p_{Y|X}(y|\text{left}) \\
 &= -\frac{3}{25} \log_2 \frac{3}{25} - \frac{22}{25} \log_2 \frac{22}{25} \approx \boxed{0.5294} \\
 H(Y|X = \text{right}) &= - \sum_{y \in Y} p_{Y|X}(y|\text{right}) \log_2 p_{Y|X}(y|\text{right}) \\
 &= -\frac{10}{31} \log_2 \frac{10}{31} - \frac{21}{31} \log_2 \frac{21}{31} \approx \boxed{0.9072}
 \end{aligned}$$

$$\begin{aligned}
 \therefore IG(Y|X) &= H(Y) - H(Y|X) \\
 &= H(Y) - P(X = \text{left})H(Y|X = \text{left}) - P(X = \text{right})H(Y|X = \text{right}) \\
 &\approx 0.7817 - \frac{25}{56}(0.5294) - \frac{31}{56}(0.9072) \approx \boxed{0.0432}
 \end{aligned}$$



Small Tree	Balanced Tree	Big Tree
<ul style="list-style-type: none"> <u>Underfitting</u>: Simple, can't handle important distinctions in data 	<ul style="list-style-type: none"> <u>"Occam's Razor"</u>: Find the simplest hypothesis that fits observations Ideally, smaller trees with informative nodes near the root 	<ul style="list-style-type: none"> <u>Overfitting</u>: Complex, incorporates meaningless noise into decisions Difficult for humans to interpret Computationally inefficient (redundant, nonsense attributes)

Pitfalls of Decision Trees

- As you descend paths in a decision tree, training set data is further subdivided and exponentially decreases
- Greedy algorithms do not necessarily yield the global optimum

KNN	Decision Trees
<ul style="list-style-type: none"> ➤ One hyperparameter ➤ Can choose what to measure distance of, how to measure distance 	<ul style="list-style-type: none"> ➤ Faster at test-time ➤ Easier to interpret for a non-professional ➤ Easier for discrete features, missing values, poorly-scaled data

Bias-Variance Decomposition

Generalization: Ability of models to predict from unseen examples (in the same distribution as training data)

- Models can overfit and underfit simultaneously. Simple models underfit more, complex models overfit more.

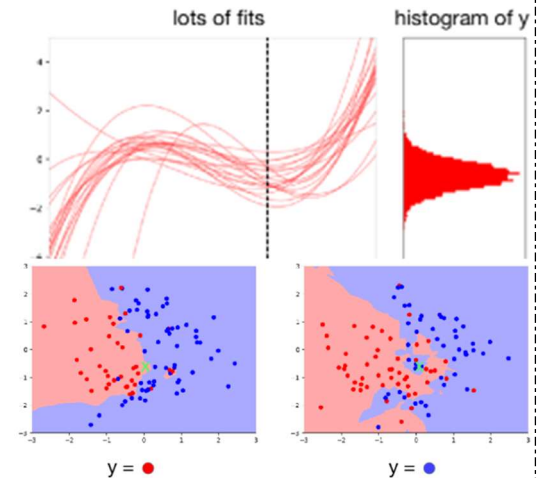
Bias-Variance Decomposition (BVD): A way to quantify underfitting and overfitting

- Data Generating Distribution ($p_{\mathbf{x},T}(\mathbf{x},t)$):** A hypothetical probability distribution function from which we assume all training/test data \mathcal{D} originates. We assume data is **independent and identically distributed (i.i.d.)**.

Computation

- Create a training set \mathcal{D} using samples from $p_{\mathbf{x},T}(\mathbf{x},t)$
- Pick a fixed query point \mathbf{x} in \mathcal{D} , find output $Y = y(\mathbf{x})$
- Sample a T value from $p_{T|X}(t|\mathbf{x})$, the hypothetical probability distribution of outputs, given an input
- Compute loss using $\mathcal{L}(Y, T)$
- By repeating steps 2-5 enough times, we find the expected loss, $\mathbb{E}[\mathcal{L}(Y, T)|X]$. We want to minimize this value to reduce bias and variance in our model.

Y and T are both random variables (RV), but when we find \mathbb{E} , we treat one of Y and T as a RV at a time and the other as a constant.



BVD of Mean Squared Error

Recall, MSE's loss function is $\mathcal{L}(Y, T) = (Y - T)^2$

Treating T as RV, Y as arbitrary output	Treating Y as RV, T as arbitrary output
$\begin{aligned}\mathbb{E}[\mathcal{L}(Y, T) X] &= \mathbb{E}[(Y - T)^2 X] \\ &= \mathbb{E}[Y^2 - 2YT + T^2 X] \\ &= \mathbb{E}[Y^2] - 2Y\mathbb{E}[T X] + \mathbb{E}[T^2 X] \\ &= \mathbb{E}[Y^2] - 2Y\mathbb{E}[T X] + \mathbb{E}[T X]^2 + \text{Var}[T X] \\ &= \mathbb{E}[(Y - \mathbb{E}[T X])^2] + \text{Var}[T X]\end{aligned}$	$\begin{aligned}\mathbb{E}[\mathcal{L}(Y, T) X] &= \mathbb{E}[(Y - \mathbb{E}[T X])^2] + \text{Var}[T X] \quad (\text{from result on the left}) \\ &= \mathbb{E}[(Y - Y^*)^2 X] + \text{Var}[T X] \\ &= Y^{*2} - 2Y^*\mathbb{E}[Y X] + \mathbb{E}[Y^2 X] + \text{Var}[T X] \\ &= Y^{*2} - 2Y^*\mathbb{E}[Y X] + \mathbb{E}[Y X]^2 + \text{Var}[Y X] + \text{Var}[T X] \\ &= \underbrace{(Y^* - \mathbb{E}[Y X])^2}_{\text{bias}} + \underbrace{\text{Var}[Y X]}_{\text{variance}} + \underbrace{\text{Var}[T X]}_{\text{Bayes error}}\end{aligned}$
To minimize $(Y - \mathbb{E}[T X])^2$, set $Y = \mathbb{E}[T X] = Y^*$; predict the expected output of a sample, the Bayes optimal decision.	
The second term is Bayes Error , the inherent unpredictability/noise of T . This is the best error algorithms can hope to get.	

Bayes Optimal: An algorithm that makes Bayes optimal decisions, ones that minimize expected loss; any error in this algorithm is due to Bayes error and it cannot be further improved.

Bias-Variance Trade-off: Idea that by tuning parameters/hyperparameters, one can increase bias but at the cost of decreasing variance and vice versa.

Bias: Model accuracy, related to <u>underfitting</u>	Model can't learn all patterns in data Model has enough data for stable estimates	High bias Low variance
Variance: Model variability, related to <u>overfitting</u>	Model learns patterns in data Model incorporates noise/quirks of sampled data	Low bias High variance

Bagging/Bootstrap Aggregation

Ensemble Method: Using the majority prediction of multiple learning algorithms to get a prediction

- Small ensembles can be better than a single good model

Bagging/Bootstrap Aggregation: A type of ensemble method where n training sets are “sampled” from $p_{X,T}(\mathbf{x}, t)$, a model is trained on each set, and each prediction $y_i(\mathbf{x})$ (for the model on the i -th set) is averaged/compared.

$$\text{Regression} \\ y(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n y_i(\mathbf{x})$$

$$\text{Classification (Binary)} \\ y(\mathbf{x}) = \mathbb{I} \left(\frac{1}{n} \sum_{i=1}^n y_i(\mathbf{x}) > 0.5 \right)$$

- In reality, $p_{X,T}(\mathbf{x}, t)$ is often expensive/impossible to sample from, and we only have one training set \mathcal{D} .
- Thus, assume data in \mathcal{D} is independent, use $p_{\mathcal{D}}(\mathbf{x}, t)$ as a proxy for $p_{X,T}(\mathbf{x}, t)$, create n datasets using resamples/bootstrap samples drawn with replacement from $p_{\mathcal{D}}(\mathbf{x}, t)$.
- As $|\mathcal{D}| \rightarrow \infty$, $p_{\mathcal{D}}(\mathbf{x}, t) \rightarrow p_{X,T}(\mathbf{x}, t)$

Bias is the same (as predictions are same)	Variance is reduced (as we average independent predictions)
$\mathbb{E}[y(\mathbf{x})] = \mathbb{E}[y_i(\mathbf{x})]$	$\text{Var}[y(\mathbf{x})] = \text{Var} \left[\frac{1}{n} \sum_{i=1}^n y_i(\mathbf{x}) \right] = \frac{1}{n^2} \sum_{i=1}^n \text{Var}[y_i(\mathbf{x})] = \frac{1}{n} \text{Var}[y_i(\mathbf{x})]$

Pitfalls of Bagging:

- Resampled data might not be independent! We thus try reducing correlation between resampled data sets by introducing variability
 - eg. Average over multiple algorithms, multiple configurations of an algorithm
 - eg. In **random forests** (ie. ensemble of decision trees), choose a subset of features for each tree, have that tree only split according to those features
- Does not reduce bias (at least in the case of squared error)
- Equal weights to ensemble members may not be ideal – weighted ensembling is often better for radically-different members


Linear Regression

Partial Derivative: Of function $f: \mathbb{R}^n \rightarrow \mathbb{R}$, the derivative with respect to (w.r.t.) a variable

- For instance, in $\frac{\partial f}{\partial x} = \frac{\partial}{\partial x} f$, differentiate f w.r.t. x , treating all other variables as constants.
- As a notational shorthand, $\frac{\partial f}{\partial \mathbf{x}} = \frac{\partial f}{\partial (x_1, \dots, x_n)} = \left[\frac{\partial f}{\partial x_1} \quad \dots \quad \frac{\partial f}{\partial x_n} \right]$
- eg. $\frac{\partial}{\partial x} (x^2 y + x y^2 + y) = 2yx + y^2$

Gradient (∇): Of function $f: \mathbb{R}^n \rightarrow \mathbb{R}$, vector $\left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right) \in \mathbb{R}^n$.

- eg. $\nabla (x^2 y + x y^2 + y) = (2yx + y^2, x^2 + 2xy + 1)$
- If $\frac{\partial}{\partial x_i} f(x) = \alpha_i$, then $\nabla f = (\alpha_1, \dots, \alpha_n) = \boldsymbol{\alpha}$
- If A is a matrix and \mathbf{x}, b are vectors, then $\frac{\partial}{\partial \mathbf{x}} (A\mathbf{x} + b) = A$



Convex Function: Functions where $f'' > 0$ or $f'' < 0$. AKA the line segment between any x_0, x_1 is above/below $f(x)$,

$$f((1-\lambda)x_0 + \lambda x_1) \leq (1-\lambda)f(x_0) + \lambda f(x_1)$$

- Recall a critical point at $f' = 0$ is either a local maximum or local minimum
- The local maximum/minimum of a convex function is the global maximum/minimum

Vectorization: The practice of turning loops in code to vector/matrix operations.

- *Simpler & readable* code – no dummy variables/indices
- *Computationally faster* – linear algebra libraries are highly-optimized, matrix multiplication is very fast, parallelizable and good for GPU, cut down on slowness of Python interpreter overhead
 - Why use Python vs. C/C++ if performance matters? Because Python's quicker to write code in, it's better for prototyping, quickly experimenting, & communicating algorithms to others. Most code is done in libraries compiled in other languages too. C/C++ is good for optimizing algorithms.
- Most “math” in of a machine learning algorithm is just matrices/vectors anyways

Linear Regression: A parametric regression model that finds the best linear relationship between two variables

- Involves many *modular components* very typical of machine learning:
 - Model* describes variable relationships
 - Loss function* quantifies how bad a fit to data is
 - Optimization Algorithm* minimizes the loss function
 - Regularizer* quantifies our preference for candidate models (*optional*)
- Goal is to learn a function y such that $\forall (\mathbf{x}, t) \in \mathcal{D}, t \approx y(\mathbf{x})$

Model: Let $\mathbf{x} \in \mathbb{R}^D$ be input, $y(\mathbf{x})$ be output, $\mathbf{w} \in \mathbb{R}^D$ be weight, $b \in \mathbb{R}$ be bias/intercept. Parameters are b and \mathbf{w} .

$$y(\mathbf{x}) = b + \sum_{d=1}^D (w_d x_d) = b + \mathbf{w} \cdot \mathbf{x}$$

Vectorization in Linear Regression

For N training set vectors of dimension D .

Store all training set vectors into a design matrix of size $\mathbb{R}^{N \times D}$	$X = \begin{bmatrix} \bar{x}_1^T \\ \vdots \\ \bar{x}_N^T \end{bmatrix} = \begin{bmatrix} (x_1)_1 & \dots & (x_1)_D \\ \vdots & \ddots & \vdots \\ (x_N)_1 & \dots & (x_N)_D \end{bmatrix}$	$X\mathbf{w} + \mathbf{b} = \begin{bmatrix} (x_1)_1 & \dots & (x_1)_D \\ \vdots & \ddots & \vdots \\ (x_N)_1 & \dots & (x_N)_D \end{bmatrix} \begin{bmatrix} w_1 \\ \vdots \\ w_D \end{bmatrix} + \begin{bmatrix} b \\ \vdots \\ b \end{bmatrix}$
Store all outputs into a target vector of size \mathbb{R}^N	$\mathbf{t} = (t_1, \dots, t_N)$	$= \begin{bmatrix} \mathbf{w} \cdot \mathbf{x}_1 + b \\ \vdots \\ \mathbf{w} \cdot \mathbf{x}_N + b \end{bmatrix}$
Store weights for each x_d as a vector of size \mathbb{R}^D	$\mathbf{w} = (w_1, \dots, w_D)$	$= \begin{bmatrix} y(\mathbf{x}_1) \\ \vdots \\ y(\mathbf{x}_N) \end{bmatrix} \approx \mathbf{t} = \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix}$
Store bias as vector in \mathbb{R}^N	$\mathbf{b} = (b, \dots, b)$	

We can simplify $X\mathbf{w} + \mathbf{b}$ into $X\mathbf{w}$ by cramming value b into X and \mathbf{w} . We could also just label b as w_0 .

$$X = \begin{bmatrix} 1 & \mathbf{x}_1^T \\ \vdots & \vdots \\ 1 & \mathbf{x}_N^T \end{bmatrix} = \begin{bmatrix} 1 & (x_1)_1 & \cdots & (x_1)_D \\ \vdots & \vdots & \ddots & \vdots \\ 1 & (x_N)_1 & \cdots & (x_N)_D \end{bmatrix}, \mathbf{w} = \begin{bmatrix} b \\ w_1 \\ \vdots \\ w_D \end{bmatrix}, \text{ then } X\mathbf{w} = \begin{bmatrix} 1 & (x_1)_1 & \cdots & (x_1)_D \\ \vdots & \vdots & \ddots & \vdots \\ 1 & (x_N)_1 & \cdots & (x_N)_D \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ \vdots \\ w_D \end{bmatrix} = \begin{bmatrix} \mathbf{w} \cdot \mathbf{x}_1 + b \\ \vdots \\ \mathbf{w} \cdot \mathbf{x}_N + b \end{bmatrix}$$

Loss/Cost Function: Use a modified squared error loss function with $\frac{1}{2}$ to simplify derivative calculations (see below).

$$\begin{aligned} \mathcal{L}(y(\mathbf{x}), t) &= \frac{1}{2} (y(\mathbf{x}) - t)^2 \\ \mathcal{J}(y(\mathbf{x}), t) &= \frac{1}{N} \sum_{n=1}^N \mathcal{L}(y(\mathbf{x}_n), t_n) \\ &= \frac{1}{2N} \sum_{n=1}^N (y(\mathbf{x}_n) - t_n)^2 \\ &= \frac{1}{2N} [(y(\mathbf{x}_1) - t_1)^2 + \cdots + (y(\mathbf{x}_N) - t_N)^2] \end{aligned} \quad \begin{aligned} &= \frac{1}{2N} \left\| \begin{bmatrix} y(\mathbf{x}_1) \\ \vdots \\ y(\mathbf{x}_N) \end{bmatrix} - \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix} \right\|^2 \\ &= \frac{1}{2N} \left\| \begin{bmatrix} \mathbf{w} \cdot \mathbf{x}_1 + b \\ \vdots \\ \mathbf{w} \cdot \mathbf{x}_N + b \end{bmatrix} - \mathbf{t} \right\|^2 \\ &= \frac{1}{2N} \|X\mathbf{w} - \mathbf{t}\|^2 \end{aligned}$$

We then write \mathcal{J} as a function of its weights and bias, $\mathcal{J}(\mathbf{w})$, to make it suitable for optimizing our parameter, \mathbf{w} .

Optimization Algorithm: Find $\text{argmin}_{\mathbf{w} \in \mathbb{R}^D} \mathcal{J}(\mathbf{w})$, the \mathbf{w} that minimizes a non-negative $\mathcal{J}(\mathbf{w})$.

- **Direct Solution:** Find critical point! Set $\nabla \mathcal{J}(\mathbf{w}) = 0$, solve for \mathbf{w} (possible in linear regression, rarely possible elsewhere)

$$\nabla \mathcal{J}(\mathbf{w}) = \frac{1}{N} (X^T X \mathbf{w} - X^T \mathbf{t}) = 0$$

$$\therefore \mathbf{w} = (X^T X)^{-1} X^T \mathbf{t}$$

- **Iterative Solution:** Repeatedly apply an update rule that gradually takes us closer to the solution
 - **Iterative Algorithm:** An algorithm that is repeatedly updates a variable until it meets criteria
 - **(Batch) Gradient Descent:** Iterative algorithm that adjusts weights in direction of steepest descent
 - **Direction of Steepest Ascent:** The direction of ∇f
 - **Direction of Steepest Descent:** The direction of $-\nabla f$

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla \mathcal{J}(\mathbf{w})$$


$$= \mathbf{w} - \frac{\alpha}{N} \sum_{n=1}^N (y(\mathbf{x}_n) - t_n) \mathbf{x}_n$$

where α = step size/learning rate, a hyperparameter.

- Values of α are usually very small (eg. 0.01)
- Need smaller α for total loss (N times as small) vs. average loss.
- Use gradient descent as it's versatile, often easier than direct solving, computationally efficient
 - eg. In linear regression, for $\mathbf{w}, \mathbf{x} \in \mathbb{R}^D$ and N samples of \mathbf{x} , the direct solution of $\mathbf{w} = (X^T X)^{-1} X^T \mathbf{t}$ requires matrix inversion, which is $\mathcal{O}(D^3)$
 - eg. Gradient descent cost is $\mathcal{O}(ND)$ and can be even less.

Feature Mapping: Using linear regression on a non-linear relationship by transforming input space using feature mapping $\psi: \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}$ into something linear, comparing transformed inputs with outputs.

- **Polynomial Feature Mapping:** Fitting data with $y(x) = \sum_{n=0}^N w_n x^n = w_0 + w_1 x + w_2 x^2 + \cdots + w_N x^N$
 - Here, we have $\psi: \mathbb{R} \rightarrow \mathbb{R}^N$ by $\psi(x) = (x^0, \dots, x^N)$, so $y(\mathbf{x}) = \mathbf{w} \cdot \psi(\mathbf{x})$.
 - N is a hyperparameter. $N = 1$ is just regular linear regression.



Optimization – Derivation of direct solution, linear regression

Since $\nabla \mathcal{J}(\mathbf{w}) = \left(\frac{\partial}{\partial w_1} \mathcal{J}(\mathbf{w}), \dots, \frac{\partial}{\partial w_D} \mathcal{J}(\mathbf{w}) \right)$, we will consider the derivative w.r.t. a component.

$$\begin{aligned}
 \frac{\partial}{\partial w_d} \mathcal{J}(\mathbf{w}) &= \frac{\partial}{\partial w_d} \left(\frac{1}{2N} \|X\mathbf{w} - \mathbf{t}\|^2 \right) \\
 &= \frac{1}{2N} \frac{\partial}{\partial w_d} \left(\left\| \begin{bmatrix} \mathbf{w} \cdot \mathbf{x}_1 + b \\ \vdots \\ \mathbf{w} \cdot \mathbf{x}_N + b \end{bmatrix} - \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix} \right\|^2 \right) \\
 &= \frac{1}{2N} \frac{\partial}{\partial w_d} \left(\sqrt{\sum_{n=1}^N [(\mathbf{w} \cdot \mathbf{x}_n + b) - t_n]^2} \right)^2 \\
 &= \frac{1}{2N} \frac{\partial}{\partial w_d} \sum_{n=1}^N [(\mathbf{w} \cdot \mathbf{x}_n + b) - t_n]^2 \\
 &= \frac{1}{2N} \sum_{n=1}^N 2[(\mathbf{w} \cdot \mathbf{x}_n + b) - t_n] \frac{\partial}{\partial w_d} [\mathbf{w} \cdot \mathbf{x}_n + b - t_n] \quad (\text{chain rule works the same}) \\
 &= \frac{1}{N} \sum_{n=1}^N [(\mathbf{w} \cdot \mathbf{x}_n + b) - t_n] \frac{\partial}{\partial w_d} [(w_1(x_n)_1 + \dots + w_d(x_n)_d + \dots + w_D(x_n)_D) + b - t_n] \\
 &= \frac{1}{N} \sum_{n=1}^N [(\mathbf{w} \cdot \mathbf{x}_n + b) - t_n] (x_n)_d \quad (= \text{iterative solution}) \\
 &= \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b)(x_n)_d - \frac{1}{N} \sum_{n=1}^N t_n (x_n)_d \\
 \therefore \nabla \mathcal{J}(\mathbf{w}) &= \begin{bmatrix} \frac{\partial}{\partial w_1} \mathcal{J}(\mathbf{w}) \\ \vdots \\ \frac{\partial}{\partial w_D} \mathcal{J}(\mathbf{w}) \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b)(x_n)_1 \\ \vdots \\ \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b)(x_n)_D \end{bmatrix} - \begin{bmatrix} \frac{1}{N} \sum_{n=1}^N t_n (x_n)_1 \\ \vdots \\ \frac{1}{N} \sum_{n=1}^N t_n (x_n)_D \end{bmatrix} = \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b) \overrightarrow{x_n} - \frac{1}{N} \sum_{n=1}^N t_n \overrightarrow{x_n}
 \end{aligned}$$

We then vectorize by turning all summations into matrices. Make sure to check dimensions for matrix multiplication!

$$\begin{aligned}
 \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b) \overrightarrow{x_n} &= \frac{1}{N} \begin{bmatrix} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b)(x_n)_1 \\ \vdots \\ \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b)(x_n)_D \end{bmatrix} \left(\text{expand outer vector } \mathbf{x}_n = \begin{bmatrix} (x_n)_1 \\ \vdots \\ (x_n)_D \end{bmatrix} \right) \\
 &= \frac{1}{N} \begin{bmatrix} (\mathbf{w} \cdot \mathbf{x}_1 + b)(x_1)_1 + \dots + (\mathbf{w} \cdot \mathbf{x}_N + b)(x_N)_1 \\ \vdots \\ (\mathbf{w} \cdot \mathbf{x}_1 + b)(x_1)_D + \dots + (\mathbf{w} \cdot \mathbf{x}_N + b)(x_N)_D \end{bmatrix} \quad (\text{expand summation}) \\
 &= \frac{1}{N} \begin{bmatrix} (x_1)_1 & \dots & (x_N)_1 \\ \vdots & \ddots & \vdots \\ (x_1)_D & \dots & (x_N)_D \end{bmatrix} \begin{bmatrix} \mathbf{w} \cdot \mathbf{x}_1 + b \\ \vdots \\ \mathbf{w} \cdot \mathbf{x}_N + b \end{bmatrix} \quad (\text{factor into matrices}) \\
 &= \frac{1}{N} [\mathbf{x}_1 \quad \dots \quad \mathbf{x}_N] X \mathbf{w} \\
 &= \frac{1}{N} X^T X \mathbf{w} \\
 \frac{1}{N} \sum_{n=1}^N t_n \mathbf{x}_n &= \frac{1}{N} \begin{bmatrix} \sum_{n=1}^N t_n (x_n)_1 \\ \vdots \\ \sum_{n=1}^N t_n (x_n)_D \end{bmatrix} = \frac{1}{N} \begin{bmatrix} t_1(x_1)_1 + \dots + t_N(x_N)_1 \\ \vdots \\ t_1(x_1)_D + \dots + t_N(x_N)_D \end{bmatrix} = \frac{1}{N} \begin{bmatrix} (x_1)_1 & \dots & (x_N)_D \\ \vdots & \ddots & \vdots \\ (x_1)_D & \dots & (x_N)_D \end{bmatrix} \begin{bmatrix} t_1 \\ \vdots \\ t_N \end{bmatrix} = \frac{1}{N} X^T \mathbf{t}
 \end{aligned}$$

From which you can conclude

$$\begin{aligned}\therefore \nabla \mathcal{J}(\mathbf{w}) &= \frac{1}{N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n + b) \mathbf{x}_n - \frac{1}{N} \sum_{n=1}^N t_n \mathbf{x}_n \\ &= \frac{1}{N} (X^T X \mathbf{w} - X^T \mathbf{t}) = 0 \\ X^T X \mathbf{w} &= X^T \mathbf{t} \\ \therefore \mathbf{w} &= (X^T X)^{-1} X^T \mathbf{t}\end{aligned}$$

For future reference, here're useful formulas, which you can find more of [here](#) and [here](#):

Squared norm to vector dot product	$\ \mathbf{x}\ ^2 = \mathbf{x}^T \cdot \mathbf{x} = \mathbf{x} \cdot \mathbf{x}^T$
Derivative product of functions of x	$\frac{\partial}{\partial \mathbf{x}} f(\mathbf{x})^T g(\mathbf{x}) = f(\mathbf{x})^T \frac{\partial g}{\partial \mathbf{x}} + g(\mathbf{x})^T \frac{\partial f}{\partial \mathbf{x}}$ for $f, g: \mathbb{R}^n \rightarrow \mathbb{R}^n$
Quadratic form	$\mathbf{x}^T A \mathbf{x}$ where A is a symmetric matrix (ie. $A_{i,j} = A_{j,i}$). $\frac{\partial}{\partial \mathbf{x}} \mathbf{x}^T A \mathbf{x} = 2 \mathbf{x}^T A$
Matrix product to summation	$(XY)_{i,k} = \sum_{j=1}^n X_{i,j} Y_{j,k}$

Here's another direct solution. Note some differing notation, and that I left out b because leaving it in is annoying.

$$\begin{aligned}\mathcal{J}(\mathbf{w}) &= \frac{1}{2N} \sum_{n=1}^N (\mathbf{w} \cdot \mathbf{x}_n - t_n)^2 \\ &= \frac{1}{2N} \sum_{n=1}^N \left(\sum_{d=1}^D (w_d(x_n)_d) - t_n \right)^2 \\ \frac{\partial}{\partial w_{d^*}} \mathcal{J}(\mathbf{w}) &= \frac{1}{2N} \sum_{n=1}^N \left[2 \left(\sum_{d=1}^D (w_d(x_n)_d) - t_n \right) \frac{\partial}{\partial w_{d^*}} \left(\sum_{d=1}^D (w_d(x_n)_d) - t_n \right) \right] \quad (\text{where } d^* \in \{1, \dots, D\}) \\ &= \frac{1}{N} \sum_{n=1}^N \left(\sum_{d=1}^D (w_d(x_n)_d) - t_n \right) (x_n)_{d^*} \quad \left(= \frac{1}{N} \sum_{n=1}^N (\vec{w} \cdot \vec{x}_n - t_n) (x_n)_{d^*}, \text{iterative solution} \right) \\ &= \frac{1}{N} \sum_{n=1}^N \left(\sum_{d=1}^D w_d(x_n)_d \right) (x_n)_{d^*} - \frac{1}{N} \sum_{n=1}^N t_n (x_n)_{d^*} \\ &= \frac{1}{N} \sum_{d=1}^D \left(\sum_{n=1}^N (x_{d^*}^T)_n (x_n)_d \right) w_d - \frac{1}{N} \sum_{n=1}^N (x_{d^*}^T)_n t_n \\ &= \sum_{d=1}^D A_{d^*,d} w_d - \mathbf{c}_{d^*} \\ A_{d^*,d} &= \frac{1}{N} \sum_{n=1}^N (x_{d^*}^T)_n (x_n)_d = \frac{1}{N} \sum_{n=1}^N X_{d^*,n}^T X_{n,d} \\ \mathbf{c}_{d^*} &= \frac{1}{N} \sum_{n=1}^N (x_{d^*}^T)_n t_n = \frac{1}{N} \sum_{n=1}^N X_{d^*,n}^T t_n \\ &\Rightarrow \begin{aligned} A &= \frac{1}{N} X^T X \\ \mathbf{c} &= \frac{1}{N} X^T \mathbf{t} \end{aligned}\end{aligned}$$

From which you can conclude

$$\begin{aligned}\therefore \nabla \mathcal{J}(\mathbf{w}) &= \begin{bmatrix} \sum_{d=1}^D A_{1,d} w_d \\ \vdots \\ \sum_{d=1}^D A_{N,d} w_d \end{bmatrix} - \begin{bmatrix} c_1 \\ \vdots \\ c_D \end{bmatrix} \\ &= \begin{bmatrix} A_{1,1} w_1 + \dots + A_{1,D} w_D \\ \vdots \\ A_{N,1} w_1 + \dots + A_{N,D} w_D \end{bmatrix} - \mathbf{c} \\ &= \begin{bmatrix} A_{1,1} & \dots & A_{1,D} \\ \vdots & \ddots & \vdots \\ A_{N,1} & \dots & A_{N,D} \end{bmatrix} \begin{bmatrix} w_1 \\ \vdots \\ w_D \end{bmatrix} - \mathbf{c} \\ &= A \mathbf{w} - \mathbf{c} = 0 \\ \therefore \mathbf{w} &= A^{-1} \mathbf{c} \\ &= \frac{1}{N} (X^T X)^{-1} X^T \mathbf{t}\end{aligned}$$

You can also obtain the solution by working directly with matrix derivatives if you're comfortable.

$$\begin{aligned}
 \mathcal{J}(\mathbf{w}) &= \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|^2 \\
 &= \frac{1}{2N} (\mathbf{X}\mathbf{w} - \mathbf{t})^T (\mathbf{X}\mathbf{w} - \mathbf{t}) \quad \text{By } \|\mathbf{x}\|^2 = \mathbf{x}^T \cdot \mathbf{x} = \mathbf{x} \cdot \mathbf{x}^T \\
 \frac{\partial}{\partial \mathbf{w}} \mathcal{J}(\mathbf{w}) &= \frac{1}{2N} \left[2(\mathbf{X}\mathbf{w} - \mathbf{t})^T \frac{\partial}{\partial \mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{t}) \right] \quad \text{By } \frac{\partial}{\partial \mathbf{x}} f(\mathbf{x})^T g(\mathbf{x}) = f(\mathbf{x})^T \frac{\partial g}{\partial \mathbf{x}} + g(\mathbf{x})^T \frac{\partial f}{\partial \mathbf{x}} \text{ (where } f(\mathbf{w}) = g(\mathbf{w}) = \mathbf{X}\mathbf{w} - \mathbf{t}) \\
 &= \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{t})^T \mathbf{X} \quad \text{By } \frac{\partial}{\partial \mathbf{x}} (A\mathbf{x} + \mathbf{b}) = A \\
 &= \frac{1}{N} [(\mathbf{X}\mathbf{w})^T - \mathbf{t}^T] \mathbf{X} \quad \text{Since } \mathbf{X}\mathbf{w} \text{ and } \mathbf{t} \text{ are } n \times 1 \text{ vectors, } (\mathbf{X}\mathbf{w} - \mathbf{t})^T = (\mathbf{X}\mathbf{w})^T - \mathbf{t}^T \\
 &= \frac{1}{N} (\mathbf{w}^T \mathbf{X}^T \mathbf{X} - \mathbf{t}^T \mathbf{X}) = 0 \\
 &\quad \mathbf{w}^T \mathbf{X}^T \mathbf{X} = \mathbf{t}^T \mathbf{X} \\
 &\quad \mathbf{w}^T = \mathbf{t}^T \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \\
 &\quad \therefore \mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}
 \end{aligned}$$

Regularization: A function \mathcal{R} artificially added to a loss function to penalize unwanted parameter values.

- **L^2/l_2 -Regularization/Penalty:** Penalizes high magnitudes of \mathbf{w} . The $\frac{1}{2}$ simplifies derivative calculations.

$$\mathcal{R}(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2$$

- **Ridge Regression:** Least squares error combined with L^2 regularization, which punishes both bad fits (\mathcal{J}) and high weights (\mathcal{R}).

$$\begin{aligned}
 \mathcal{J}_{\text{reg}}(\mathbf{w}) &= \mathcal{J}(\mathbf{w}) + \lambda \mathcal{R}(\mathbf{w}) \\
 &= \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2
 \end{aligned}$$


- λ is a hyperparameter that controls the penalization of large weights.
- We ignore $\frac{1}{N}$ in the cost function for simplicity.

Direct Solution:

$$\begin{aligned}
 \nabla \mathcal{J}_{\text{reg}}(\mathbf{w}) &= \nabla \mathcal{J}(\mathbf{w}) + \nabla \lambda \mathcal{R}(\mathbf{w}) \\
 &= \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{t} + \lambda \mathbf{w} = 0 \\
 \therefore \mathbf{w} &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{t}
 \end{aligned}$$

Iterative Solution:

$$\begin{aligned}
 \mathbf{w} &\leftarrow \mathbf{w} - \alpha \nabla \mathcal{J}_{\text{reg}}(\mathbf{w}) \\
 &= \mathbf{w} - \alpha (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{t} + \lambda \mathbf{w}) \\
 &= \mathbf{w} (1 - \alpha \lambda) - \alpha (\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{t})
 \end{aligned}$$



Small α	Balanced α	High α
➤ Slow progress, convergence takes a long time	➤ Optimal values usually 0.001 to 0.1 ➤ Training curve (above) can diagnose issues	➤ Instability, oscillations, misses optimum


Stochastic Gradient Descent: Gradient descent that chooses a random \mathbf{x} of \mathcal{D} to calculate loss, instead of all of \mathcal{D} .

- Instead of $\mathcal{J}(\mathbf{w}) = \frac{1}{2N} \|\mathbf{X}\mathbf{w} - \mathbf{t}\|^2$, we have $\mathcal{J}(\mathbf{w}, \mathbf{x}) = \frac{1}{2} \|\mathbf{x} \cdot \mathbf{w} - t\|^2$
- Running-time of each update is independent of size of training data (ie. N). Faster!
- Mathematically, it is unbiased and equivalent to gradient descent, given sampling is random

$$\mathbb{E} \left[\frac{\partial \mathcal{J}_{\text{stochastic}}(\mathbf{w}, \mathbf{x})}{\partial \mathbf{w}} \right] = \frac{1}{N} \sum_{n=1}^N \frac{\partial \mathcal{J}_{\text{stochastic}}(\mathbf{w}, \mathbf{x}_n)}{\partial \mathbf{w}} = \frac{\partial \mathcal{J}_{\text{batch}}(\mathbf{w})}{\partial \mathbf{w}}$$

Mini-Batch Gradient Descent: Similar to stochastic gradient descent, but chooses a subset $\{\mathbf{x}_1, \dots, \mathbf{x}_k\}$ of \mathcal{D} .

- Batch/subset size is a hyperparameter.
- In theory, get i.i.d. examples from \mathcal{D} . In practice, we permute training set, go through it sequentially. Each pass over all data is called an **epoch**.



Small Batch Size	Balanced Batch Size	High Batch Size
<ul style="list-style-type: none"> ➤ Faster updates ➤ High variance in $\mathcal{J}(\mathbf{w}, \mathbf{x})$, noisier movement ➤ Cannot exploit speed of vectorization 	<ul style="list-style-type: none"> ➤ Batch size is hyperparameter ➤ A reasonable size might be 100 	<ul style="list-style-type: none"> ➤ Slower updates ➤ Less variance in $\mathcal{J}(\mathbf{w}, \mathbf{x})$, more accurate movement ➤ Exploits speed of vectorization

**Learning rate affects how batch sizes affect noise, so when testing, it's best to start with high learning rates and move down*

Linear Classification

Binary Linear Classification: A parametric classification model that, given a $\mathbf{x} \in \mathbb{R}^d$, predicts a binary target t

- Depending on what's convenient, use $t \in \{0,1\}$ or $t \in \{-1,1\}$ for negative/positive examples
- The parameter threshold r splits input space into two

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b, \quad y(\vec{x}) = \begin{cases} 1 & \text{if } f(\mathbf{x}) \geq r \\ 0 & \text{if } f(\mathbf{x}) < r \end{cases}$$

In practice, we can remove b and r by introducing $w_0 = b - r$, $x_0 = 1$, and squishing it into \mathbf{w} and \mathbf{x} like this:

$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} = \begin{bmatrix} b - r \\ w_1 \\ \vdots \\ w_d \end{bmatrix}, \mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}, f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} = (w_1 x_1 + \dots + w_d x_d) + b - r, y(\mathbf{x}) = \begin{cases} 1 & \text{if } f(\mathbf{x}) \geq 0 \\ 0 & \text{if } f(\mathbf{x}) < 0 \end{cases}$$

$|f(\mathbf{x})|$ is proportional to the algorithm's **confidence** in the prediction.


Input/Data Space: A subset of \mathbb{R}^d ; all possible of \mathbf{x} .

Weight Space: A subset of \mathbb{R}^d ; all possible values of \mathbf{w} .

Half-Space: A half of \mathbb{R}^n divided by a hyperplane (a plane in \mathbb{R}^{n-1}). Can be written as $w_0 x_0 + \dots + w_d x_d \geq r$


NOT Operator

\mathbf{x}^T			t	As $f(\mathbf{x}) = w_0 x_0 + w_1 x_1$, then for $y(\mathbf{x}) = t$ to hold, we have $f(\mathbf{x}_1) = w_0(1) + w_1(0) = w_0 \geq 0$ $f(\mathbf{x}_2) = w_0(1) + w_1(1) = w_0 + w_1 < 0$
x_0	x_1			
1	0		1	
1	1		0	



AND Operator

\mathbf{x}^T				t	As $f(\mathbf{x}) = w_0 x_0 + w_1 x_1 + w_2 x_2$, then for $y(\mathbf{x}) = t$ to hold, we must have $f(\mathbf{x}_1) = w_0(1) + w_1(0) + w_2(0) = w_0 < 0$ $f(\mathbf{x}_2) = w_0(1) + w_1(0) + w_2(1) = w_0 + w_2 < 0$ $f(\mathbf{x}_3) = w_0(1) + w_1(1) + w_2(0) = w_0 + w_1 < 0$ $f(\mathbf{x}_4) = w_0(1) + w_1(1) + w_2(1) = w_0 + w_1 + w_2 \geq 0$
x_0	x_1	x_2			
1	0	0		0	
1	0	1		0	
1	1	0		0	
1	1	1		1	



Feasible Region: A region of weight space that correctly classifies all \mathbf{x}_i

- Each \mathbf{x}_i adds more constraints on w_i , AKA creates more half-spaces whose intersections w_i must lie in

Feasible: The classification problem, if the feasible region is non-empty

Convex Set: A set S such that $\forall x_0, x_1 \in S$, the line segment connecting x_0, x_1 is in S . That is,

$$\lambda x_0 + (1 - \lambda) x_1 \quad (\text{where } 0 \leq \lambda \leq 1)$$

- Weighted averages/convex combinations are in S : $\lambda_1 x_1 + \dots + \lambda_n x_n$ (where $\lambda_i > 0, \sum_{i=1}^n \lambda_i = 1$)
- Both $+/-$ regions of input space are half-spaces, which are convex
- Feasible regions of weight space are intersections of convex half-spaces, which are convex

Linearly Separable: Training data, if linear decision rules can perfectly separate it. Almost never the case.

- If linearly separable, learn weights with linear programming/optimization (APM236), perceptron algorithm
- If not linearly separable, find weights that minimize average loss!
- Proving something is not linearly separable: Suppose it is linearly separable, then the half-spaces containing positive/negative points are convex, and so are the line segments connecting any two points in a half-space. Show a line between two points intersects a point in a different half-space, creating a contradiction

Strategy	Description	Plot
0-1 Loss	$\mathcal{L}(y(\mathbf{x}), t) = \mathbb{I}[y(\mathbf{x}) \neq t]$ $\mathcal{J}(y(\mathbf{x}), t) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}[y(\mathbf{x}_i) \neq t_i]$ <p>Also called the <u>miscalculation rate</u>.</p> <p><i>Problem:</i> Derivative is 0 almost everywhere and undefined at $f(\mathbf{x}) = 0$, so we can't minimize loss with gradient descent</p>	
Squared Loss	$\mathcal{L}(y(\mathbf{x}), t) = \frac{1}{2} (f(\mathbf{x}) - t)^2$ <p>Unlike 0-1 Loss, we're working with $f(\mathbf{x})$, whose codomain is \mathbb{R} so r matters more. Set $r = \frac{1}{2}$ as the in-between of 0 and 1.</p> <p><i>Problem:</i> High-confidence predictions (ie. $f(\mathbf{x})$ is high) incur large loss, even if they're correct predictions.</p>	
Logistic Function, Squared Loss	$\sigma(x) = \frac{1}{1 + e^{-x}} \in [0,1]$ $y(\mathbf{x}) = \sigma(f(\mathbf{x})) = \frac{1}{1 + e^{-f(\mathbf{x})}}$ $\mathcal{L}(y(\mathbf{x}), t) = \frac{1}{2} (y(\mathbf{x}) - t)^2$ <p>σ is the logistic function, a non-linear sigmoid (s-shaped).</p> <p><i>Problem:</i> Confidently wrong answers have low gradients, resulting in slow updates (ie. the gradient signal is weak)</p>	
Logistic Function, Cross-Entropy / Logistic Loss	$y(\mathbf{x}) = \sigma(f(\mathbf{x})) = \frac{1}{1 + e^{-f(\mathbf{x})}}$ $\mathcal{L}_{\text{CE}}(y(\mathbf{x}), t) = \begin{cases} -\ln y(\mathbf{x}) & \text{if } t = 1 \\ -\ln(1 - y(\mathbf{x})) & \text{if } t = 0 \end{cases}$ $= -t \ln y(\mathbf{x}) - (1 - t) \ln(1 - y(\mathbf{x}))$ <p>If $f(\mathbf{x}) \ll 0$, $y(\mathbf{x}) \approx 0$ and is numerically 0, too small to store, causing computational instability (CSC336). So, expand $y(\mathbf{x})$:</p> $\mathcal{L}(y(\mathbf{x}), t) = t \ln(1 + e^{-f(\mathbf{x})}) + (1 - t) \ln(1 + e^{-f(\mathbf{x})})$	

Logistic Regression: Linear classification using logistic function and cross-entropy loss

- Why is it called "regression" if we're doing classification? It's a product of history where there were no naming conventions.

Logistic function + cross-entropy loss is optimal for linear binary classification

- Smooth, continuous, differentiable
- Small loss for correct predictions, penalizes extreme misclassifications

Direct Solution:	Iterative Solution:	$\vec{w} \leftarrow \vec{w} - \alpha \nabla \mathcal{L}(y(\vec{x}), t)$
DNE, since $\sigma(x)$ is non-linear	Possible, cross-entropy loss is convex in \vec{w}	$= \vec{w} - \alpha(y(\vec{x}) - t)\vec{x}$

Linear/logistic regression are **generalized linear models** with the same gradient update rules.

Linear Regression	Logistic Regression
$\mathbf{w} \leftarrow \mathbf{w} - \nabla \mathcal{J}(\mathbf{w})$ $= \mathbf{w} - \frac{\alpha}{N} \sum_{n=1}^N (y(\mathbf{x}_n) - t_n) \mathbf{x}_n$ <p style="text-align: center;">\downarrow $y(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$</p>	$\vec{w} \leftarrow \vec{w} - \alpha \nabla \mathcal{L}(y(\mathbf{x}), t)$ $= \vec{w} - \alpha(y(\mathbf{x}) - t)\vec{x}$ <p style="text-align: center;">\downarrow $y(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x})$</p>
Note: This cost function incorporates all \mathbf{x}_i .	Note: This loss function uses a single \mathbf{x}_i .

Optimization – Derivation of iterative solution, logistic regression

$$\begin{aligned}\nabla \mathcal{L} &= \frac{\partial \mathcal{L}_{\text{CE}}}{\partial w_i} \\ &= \frac{\partial \mathcal{L}_{\text{CE}}}{\partial y} \cdot \frac{\partial y}{\partial f} \cdot \frac{\partial f}{\partial \bar{w}} \\ &= \frac{y(\mathbf{x}) - t}{y(\mathbf{x})(1 - y(\mathbf{x}))} \cdot (1 - y(\mathbf{x}))y(\mathbf{x}) \cdot \mathbf{x} \\ &= (y(\mathbf{x}) - t)\mathbf{x}\end{aligned}$$

$$\begin{aligned}\frac{\partial \mathcal{L}_{\text{CE}}}{\partial y} &= \frac{\partial}{\partial y} [-t \ln y(\mathbf{x}) - (1 - t) \ln(1 - y(\mathbf{x}))] \\ &= -\frac{t}{y(\mathbf{x})} - (1 - t) \left(-\frac{1}{1 - y(\mathbf{x})} \right) \\ &= -\frac{t}{y(\mathbf{x})} + \frac{1 - t}{1 - y(\mathbf{x})} \\ &= \frac{-t(1 - y(\mathbf{x})) + (1 - t)y(\mathbf{x})}{y(\mathbf{x})(1 - y(\mathbf{x}))} \\ &= \frac{y(\mathbf{x}) - t}{y(\mathbf{x})(1 - y(\mathbf{x}))}\end{aligned}$$

$$\frac{\partial f}{\partial \mathbf{w}} = \frac{\partial}{\partial \mathbf{w}} [\mathbf{w} \cdot \mathbf{x}] = \mathbf{x}$$

$$\frac{\partial y}{\partial f} = \frac{\partial}{\partial f} \left[\frac{1}{1 + e^{-f(\mathbf{x})}} \right] = \frac{e^{-f(\mathbf{x})}}{(1 + e^{-f(\mathbf{x})})^2} = \frac{1 + e^{-f(\mathbf{x})} - 1}{1 + e^{-f(\mathbf{x})}} \cdot \frac{1}{1 + e^{-f(\mathbf{x})}} = \left(1 - \frac{1}{1 + e^{-f(\mathbf{x})}} \right) y(\mathbf{x}) = (1 - y(\mathbf{x}))y(\mathbf{x})$$

Multi-Class Linear Classification: Given a $\mathbf{x} \in \mathbb{R}^D$, predicts a target $t \in \{1, \dots, K\}$

- **Integer Encoding:** Encoding a target as a number
 - eg. Apple is 1, orange is 2, banana is 3, $t \in \{1, 2, 3\} \in \mathbb{R}$
- **One-hot Vector/One-of- k Encoding:** Encoding a target as a unit vector in its own dimension
 - eg. Apple is (1,0,0), orange is (0,1,0), banana is (0,0,1), and $\mathbf{t} \in \mathbb{R}^3$
 - As calculations usually imply “in-between values” and use “distances” between targets, we use one-hot vectors such that these “distances” between all targets are the same.

$$f(\mathbf{x}) = W\mathbf{x} + \mathbf{b} = \begin{bmatrix} (w_1)_1 & \cdots & (w_1)_D \\ \vdots & \ddots & \vdots \\ (w_K)_1 & \cdots & (w_K)_D \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_D \end{bmatrix} + \begin{bmatrix} b \\ \vdots \\ b \end{bmatrix}, \quad y(\mathbf{x}) = \text{softmax}(f(\mathbf{x})), \quad \mathbf{t} = \begin{bmatrix} t_1 \\ \vdots \\ t_K \end{bmatrix}$$

Softmax Function: A multivariate generalization of the logistic function, $\mathbb{R}^D \rightarrow (0,1)^D$.

- Inputs are called **logits**.
- Outputs are thought of as the probability/certainty of an option. All components add up to one.

$$\text{softmax}(\mathbf{x}) = \left(\frac{e^{x_1}}{\sum_{d=1}^D e^{x_d}}, \dots, \frac{e^{x_D}}{\sum_{d=1}^D e^{x_d}} \right)$$

- Equivalent to $\sigma(x)$ when $D = 2$:

$$\text{softmax}_1(f(x)) = \frac{e^{f_1(x)}}{e^{f_1(x)} + e^{f_2(x)}} = \frac{1}{1 + e^{f_2(x) - f_1(x)}} = \frac{1}{1 + e^{-f^*(x)}} = \sigma(f^*(x))$$

Softmax function + cross-entropy loss is optimal for linear multiclass classification

Iterative Solution:

$$\begin{aligned}\mathcal{L}(y(\mathbf{x}), \mathbf{t}) &= - \sum_{k=1}^K t_k \ln y_k(\mathbf{x}) \\ &= -\mathbf{t} \cdot (\ln y_1(\mathbf{x}), \dots, \ln y_K(\mathbf{x})) \\ \mathbf{w} &\leftarrow \mathbf{w} - \alpha \nabla \mathcal{L}(y(\mathbf{x}), \mathbf{t}) \\ &= \mathbf{w} - \alpha (y(\mathbf{x}) - \tilde{\mathbf{t}})\mathbf{x} \\ &\quad \downarrow \\ y(\mathbf{x}) &= \text{softmax}(\mathbf{w} \cdot \mathbf{x})\end{aligned}$$

$$\text{Recall } \mathcal{L}_{\text{CE}}(y(\mathbf{x}), t) = \begin{cases} -\ln y(\mathbf{x}) & \text{if } t = 1 \\ -\ln(1 - y(\mathbf{x})) & \text{if } t = 0 \end{cases}$$

This is cross-entropy re-written specifically for the binary case. For multi-class, cross-entropy is $-\ln y(\mathbf{x})$, the component where $t_k = 1$.

Why is the loss $-\sum_{k=1}^K t_k \ln y_k(\mathbf{x})$? Recall \mathbf{t} is a one-hot vector.

When $t_k = 0$, then $-t_k \ln y_k(\mathbf{x}) = 0$

When $t_k = 1$, then $-t_k \ln y_k(\mathbf{x}) = -\ln y_k(\mathbf{x})$

So the summation only computes CE at the component where $t_k = 1$

Optimization – Derivation of Iterative Solution, Logistic Regression for Multiclass Linear Classification

$$\begin{aligned}
\frac{\partial \mathcal{L}(y(\vec{x}), \vec{t})}{\partial \vec{w}_k} &= \frac{\partial \mathcal{L}(y(\vec{x}), \vec{t})}{\partial f(\vec{x})} \cdot \frac{\partial f(\vec{x})}{\partial \vec{w}_k} \\
&= \frac{\partial}{\partial f(\vec{x})} \left(-\vec{t} \cdot \begin{bmatrix} \ln y_1(\vec{x}) \\ \vdots \\ \ln y_K(\vec{x}) \end{bmatrix} \right) \cdot \frac{\partial}{\partial \vec{w}_k} (W\vec{x} + \vec{b}) \\
&= \\
W\vec{x} + \vec{b} &= \begin{bmatrix} (w_1)_1 & \cdots & (w_1)_D \\ \vdots & \ddots & \vdots \\ (w_K)_1 & \cdots & (w_K)_D \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_D \end{bmatrix} + \begin{bmatrix} b_1 \\ \vdots \\ b_K \end{bmatrix} = \begin{bmatrix} \vec{w}_1 \cdot \vec{x} + b_1 \\ \vdots \\ \vec{w}_K \cdot \vec{x} + b_K \end{bmatrix} \\
\therefore \frac{\partial}{\partial \vec{w}_k} (W\vec{x} + \vec{b}) &= \frac{\partial}{\partial \vec{w}_k} \left(\begin{bmatrix} \vec{w}_1 \cdot \vec{x} + b_1 \\ \vdots \\ \vec{w}_k \cdot \vec{x} + b_k \\ \vdots \\ \vec{w}_K \cdot \vec{x} + b_K \end{bmatrix} \right) = \vec{x} \\
-\vec{t} \cdot \begin{bmatrix} \ln y_1(\vec{x}) \\ \vdots \\ \ln y_K(\vec{x}) \end{bmatrix} &= - \begin{bmatrix} t_1 \\ \vdots \\ t_K \end{bmatrix} \cdot \begin{bmatrix} \ln y_1(\vec{x}) \\ \vdots \\ \ln y_K(\vec{x}) \end{bmatrix} \\
&= - \begin{bmatrix} t_1 \\ \vdots \\ t_K \end{bmatrix} \cdot \begin{bmatrix} \ln (\text{softmax}_1(f(\vec{x}))) \\ \vdots \\ \ln (\text{softmax}_K(f(\vec{x}))) \end{bmatrix} \\
&= -t_1 \ln (\text{softmax}_1(f(\vec{x}))) - \cdots - t_K \ln (\text{softmax}_K(f(\vec{x}))) \\
&= -t_1 \ln \left(\frac{e^{f_1(\vec{x})}}{\sum_{k=1}^K e^{f_k(\vec{x})}} \right) - \cdots - t_K \ln \left(\frac{e^{f_K(\vec{x})}}{\sum_{k=1}^K e^{f_k(\vec{x})}} \right) \\
&= -t_1 (f_1(\vec{x}) - f_1(\vec{x}) \times \dots \times f_K(\vec{x})) - t_K (f_K(\vec{x}) - f_1(\vec{x}) \times \dots \times f_K(\vec{x})) \\
&= (f_1(\vec{x}) \times \dots \times f_K(\vec{x})) (t_1 f_1(\vec{x}) + \dots + t_K f_K(\vec{x})) \\
&= (f_1(\vec{x}) \times \dots \times f_K(\vec{x})) (\vec{t} \cdot f(\vec{x})) \\
\frac{\partial}{\partial f_k(\vec{x})} \left(-\vec{t} \cdot \begin{bmatrix} \ln y_1(\vec{x}) \\ \vdots \\ \ln y_K(\vec{x}) \end{bmatrix} \right) &= \frac{\partial}{\partial f_k(\vec{x})} [(f_1(\vec{x}) \times \dots \times f_K(\vec{x})) (\vec{t} \cdot f(\vec{x}))] \\
&= \frac{\partial}{\partial f_k(\vec{x})} [(f_k(\vec{x}))^2 (f_1(\vec{x}) \times \dots \times f_{k-1}(\vec{x}) \times f_{k+1}(\vec{x}) \times \dots \times f_K(\vec{x}))] \\
&= 2f_k(\vec{x}) (f_1(\vec{x}) \times \dots \times f_{k-1}(\vec{x}) \times f_{k+1}(\vec{x}) \times \dots \times f_K(\vec{x})) \\
&= 2f_1(\vec{x}) \times \dots \times f_K(\vec{x}) \\
\therefore \frac{\partial}{\partial f(\vec{x})} \left(-\vec{t} \cdot \begin{bmatrix} \ln y_1(\vec{x}) \\ \vdots \\ \ln y_K(\vec{x}) \end{bmatrix} \right) &= 2 \begin{bmatrix} f_1(\vec{x}) \times \dots \times f_K(\vec{x}) \\ \vdots \\ f_1(\vec{x}) \times \dots \times f_K(\vec{x}) \end{bmatrix} \\
&\text{ANSWER} = (y_k(\vec{x}) - t_k) \vec{x}
\end{aligned}$$


AAAAAAA UNFINISHED

$$\begin{aligned}
&\frac{e^{x_1}}{\sum_{d=1}^D e^{x_d}} \\
\ln \left(\frac{e^{x_i}}{\sum_{d=1}^D e^{x_d}} \right) &= \ln e^{x_i} - \ln(e^{x_1} + \dots + e^{x_D}) = x_i - \ln(e^{x_1}) \dots \ln(e^{x_D}) = x_i - x_1 \dots x_D \\
w_i &\leftarrow w_i - \frac{\alpha}{N} \sum_{n=1}^N ((y_n)_i - (t_n)_i) x_i
\end{aligned}$$

Gradient Checking: Process of testing correctness of derivative implementation.

- **Finite Differences:** In numerical analysis, a way to approximate derivatives with something of form $f(x + b) - f(x + a)$. Use limit:

$$\frac{\partial}{\partial x_i} f(\mathbf{x}) = \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_i + h, \dots, x_n) - f(x_1, \dots, x_i - h, \dots, x_n)}{2h}$$
 - Use **double precision floats** (preserve 15-17 significant digits)
 - Plug in a **tiny h** (eg. 10^{-10}) for good approximations
- Let a = finite differences estimate, b = your implementation. Find **relative error** = $\frac{|a-b|}{|a|+|b|}$, which should be small (eg. 10^{-6})
- Algorithms often appear to work even when math is wrong.



KNN	Linear Classifiers
➤ High variance if k is small	➤ High bias, may underfit
➤ No parameters	➤ Parameters \vec{w}, b
➤ No training, slow at test-time	➤ Needs to learn a model, fast at test-time
➤ Bad in high-dimensions, scale-sensitive	➤ Good in high dimensions, not scale-sensitive

Classification Metrics


Accuracy: Fraction of examples correctly classified, equivalent to average 0-1 loss/error rate/misclassification rate

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{P} + \text{N}} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

P	N	TP	TN	FP	FN
Positives	Negatives	True positives	True negatives	False positives, “type I” errors	False negatives, “type II” errors

Sensitivity: True positive rate $\text{Sens.} = \frac{\text{TP}}{\text{P}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$

Specificity: True negative rate $\text{Spec.} = \frac{\text{TN}}{\text{N}} = \frac{\text{TN}}{\text{TN} + \text{FP}}$



Receiver Operating Characteristic (ROC) Curve:

Tracks a binary classifier's specificity/sensitivity trade-off (top-left of curve is the "ideal")

Confusion Matrix: For a multiclass classifier, a $K \times K$ matrix where rows are true targets, columns are predicted targets, entries are frequencies

